

自然语言处理技术

Neural Net Fundamentals

神经网络基础

计算机科学与技术学院

杨沐昀 孙承杰

# 主要内容

- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 主要内容


- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 分类问题定义与表示

- We have a training dataset consisting of  $N$  samples  $\{x_i, y_i\}_{i=1}^N$
- $x_i$  are inputs, e.g. words (indices or vectors!), sentences, documents, etc.
  - Dimension  $d$
- $y_i$  are labels (one of  $C$  classes) we try to predict, for example:
  - classes: sentiment, named entities
  - other words
  - later: multi-word sequences

# 分类方法

- Training data:
  - Fixed 2D word vectors to classify
  - Using softmax/logistic regression
  - Linear decision boundary



Visualizations with ConvNetJS by Karpathy!

<http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>

- Traditional ML/Stats approach: assume  $x_i$  are fixed, train (i.e., set) softmax/logistic regression weights  $W^{C \times d}$  to determine a decision boundary (hyperplane) as in the picture.
- Method: For each  $x$ , predict:

$$p(y|x) = \frac{\exp(W_y \cdot x)}{\sum_{c=1}^C \exp(W_c \cdot x)}$$

# softmax分类方法

$$p(y|x) = \frac{\exp(W_y \cdot x)}{\sum_{c=1}^C \exp(W_c \cdot x)}$$

可以将预测函数分解为两个步骤:

1. 将矩阵 $W$  的第 $y$ 行与向量 $x$  相乘, 计算 $f_c$  ( $c = 1, \dots, C$ )

$$W_y \cdot x = \sum_{i=1}^d W_{yi} x_i = f_y$$

2. 应用softmax函数得到归一化概率:

$$p(y|x) = \frac{\exp(f_y)}{\sum_{c=1}^C \exp(f_c)} = \text{softmax}(f_y)$$

# 使用cross-entropy作为损失函数

- 对于每个训练样本  $(x, y)$ ，我们的目标是最大化正确类别 $y$ 的概率
- 或者说最小化正确类别 $y$ 的负对数概率  
( negative log probability )

$$-\log p(y|x) = -\log \left( \frac{\exp(f_y)}{\sum_{c=1}^C \exp(f_c)} \right)$$

# cross-entropy损失函数

- “cross entropy” 是一个信息论里的概念
- 如果正确的类别概率分布为 $p$ ，模型得到的类别概率分布为 $q$ ，则cross entropy可以被定义为：

$$H(p, q) = - \sum_{c=1}^C p(c) \log q(c)$$

- 假设正确的类别概率分布为正确类别对应的概率为1，其他类别对应的概率为0，即 $p=[0, \dots, 0, 1, 0, \dots, 0]$ 。那么，cross entropy 中唯一留下的项就是正确类别的负对数概率。



# cross-entropy损失函数

- 在整个数据集 $\{x_i, y_i\}_{i=1}^N$  上的cross-entropy损失函数可以表示为:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right)$$

- 为了表示方便, 我们将使用矩阵表示

$$f = Wx$$

来代替之前的表示  $f_y = f_y(x) = W_{y \cdot} x = \sum_{j=1}^d W_{yj} x_j$


# 传统机器学习的优化方法

- 机器学习的参数 $\theta$  一般可以表示成如下形式:

$$\theta = \begin{bmatrix} W_{.1} \\ \vdots \\ W_{.d} \end{bmatrix} = W(:) \in \mathbb{R}^{Cd}$$

- 通过下列计算过程来更新分类面


$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \nabla_{W_{.1}} \\ \vdots \\ \nabla_{W_{.d}} \end{bmatrix} \in \mathbb{R}^{Cd}$$



Visualizations with ConvNetJS by Karpathy

# 神经网络分类器 (Neural Network Classifiers)



- Softmax只能给出线性分类边界，分类能力受限



如何把这些样本分对？

# 神经网络分类器 (Neural Network Classifiers)



- Neural networks 可以表示更复杂的函数，得到非线性分类边界




# 主要内容

- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 神经计算 (Neural computation)



# 人工神经元 (artificial neuron)





# 神经元与logistic regression单元

$f$  = nonlinear activation (e.g. sigmoid),  $w$  = weights,  $b$  = bias,  $h$  = hidden,  $x$  = inputs

$$h_{w,b}(x) = f(w^T x + b)$$

$$f(z) = \frac{1}{1 + e^{-z}}$$




$w, b$  are the parameters of this neuron  
i.e., this logistic regression model




# 神经网络 (neural network)

- 一个神经网络相当于多个logistic regression单元在同时运行



# 神经网络


- 多层神经网络



损失函数会指导中间隐藏层变量的取值，以便更好地预测下一层的目标。


# 神经网络

- 多层神经网络




# 神经网络的矩阵表示

- 对于 $L_2$ ,
  - $a_1 = f(w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b_1)$
  - $a_2 = f(w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b_2)$
  - ...
- 使用矩阵表示, 可以写成
  - $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$
  - $\mathbf{a} = f(\mathbf{z})$
- 激活函数 $f$ 是应用于每个元素上的
  - $f([z_1, z_2, z_3]) = [f(z_1), f(z_2), f(z_3)]$



# 为什么需要f为非线性函数

- Example: function approximation, e.g., regression or classification
  - Without non-linearities, deep neural networks can't do anything more than a linear transform
  - Extra layers could just be compiled down into a single linear transform:  $W_1 W_2 x = Wx$
  - With more layers, they can approximate more complex functions!



# 主要内容

- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 神经网络应用例子：基于二分类的地名识别

- Example: **Not all museums in Paris are amazing** .
- Here: one true window, the one with Paris in its center and all other windows are “corrupt” in terms of not having a named entity location in their center.

**museums in Paris are amazing**


- “Corrupt” windows are easy to find and there are many: Any window whose center word isn’t specifically labeled as NER location in our corpus

**Not all museums in Paris**

# 前向计算 (Neural Network Feed-forward Computation)

- 使用一个3层神经网络来给输入的句子片段 $x$ 打分

$$score(x) = U^T a \in \mathbb{R}$$



- $s = score(\text{"museums in Paris are amazing"})$


$$s = U^T f(Wx + b)$$

$$x \in \mathbb{R}^{20 \times 1}, W \in \mathbb{R}^{8 \times 20}, U \in \mathbb{R}^{8 \times 1}$$



# 中间层的直觉解释

- 中间层学到的是输入的词向量之间的非线性交互 (non-linear interactions between the input word vectors)



- 例如: 当第一个词是 “museums” 时, “in” 位于第二个词的位置其作用会变得重要

# max-margin损失函数

- 训练目标的思想:使正例样本的分数变大, 使负例样本的分数变低(直到足够好为止)
- $s = \text{score}(\text{museums in Paris are amazing})$
- $s_c = \text{score}(\text{Not all museums in Paris})$
- 最小化

$$J = \max(0, 1 - s + s_c)$$

- $J$ 不是处处可导的, 但它是连续的 $\rightarrow$ 我们可以使用随机梯度下降 (SGD) 。

# max-margin损失函数


- 对于1个输入窗口（句子片段）

$$J = \max(0, 1 - s + s_c)$$

- 每个地名位于中心的窗口（正例）的score要比没有地名位于中心的窗口（负例）大1
- 完整的优化目标函数：对于每个正例，构造多个负例。然后对所有窗口的J求和。

# 计算score的简单神经网络

- $s = u^T h$
- $h = f(Wx + b)$
- $x$  (输入)



# 随机梯度下降

- 参数更新公式

$$\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$$

$\alpha$  为步长或者学习率 ( *step size or learning rate* )

- 如何计算  $\nabla_{\theta} J(\theta)$ 
  - 手动
  - 后向传播算法 (the backpropagation algorithm)

# 手动计算梯度

- 多变量求导
- Matrix calculus: Fully vectorized gradients
  - Much faster and more useful than non-vectorized gradients
  - But doing a non-vectorized gradient can be good practice

# 主要内容

- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 导数计算回顾

- Given a function with 1 output and 1 input

$$f(x) = x^3$$

- It's gradient (slope) is its derivative

$$\frac{df}{dx} = 3x^2$$



# 导数计算回顾

- Given a function with 1 output and n inputs

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$$

- It's gradient is a vector of partial derivatives with respect to each input

$$\frac{\partial f}{\partial \mathbf{x}} = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$$

# 雅可比矩阵 (Jacobian Matrix)

- Given a function with  $m$  outputs and  $n$  inputs

$$\mathbf{f}(\mathbf{x}) = [f_1(x_1, x_2, \dots, x_n), \dots, f_m(x_1, x_2, \dots, x_n)]$$

- It's Jacobian is an  $m \times n$  matrix of partial derivatives

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \dots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \dots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

$$\left( \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)_{ij} = \frac{\partial f_i}{\partial x_j}$$

# 链式规则 ( Chain Rule )

- For one-variable functions: multiply derivatives

$$\begin{aligned} z &= 3y \\ y &= x^2 \\ \frac{dz}{dx} &= \frac{dz}{dy} \frac{dy}{dx} = (3)(2x) = 6x \end{aligned}$$

- For multiple variables at once: multiply Jacobians

$$\begin{aligned} \mathbf{h} &= f(\mathbf{z}) \\ \mathbf{z} &= \mathbf{W}\mathbf{x} + \mathbf{b} \\ \frac{\partial \mathbf{h}}{\partial \mathbf{x}} &= \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \dots \end{aligned}$$

# 雅可比矩阵示例： Elementwise activation Function的求导

$$\mathbf{h} = f(\mathbf{z}) \quad \frac{\partial \mathbf{h}}{\partial \mathbf{z}} = ?$$
$$h_i = f(z_i)$$

- Function has n outputs and n inputs  $\rightarrow$  n by n Jacobian

$$\left( \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right)_{ij} = \frac{\partial h_i}{\partial z_j} = \frac{\partial}{\partial z_j} f(z_i) = \begin{cases} f'(z_i) & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$


$$\frac{\partial \mathbf{h}}{\partial \mathbf{z}} = \begin{pmatrix} f'(z_1) & & 0 \\ & \ddots & \\ 0 & & f'(z_n) \end{pmatrix} = \text{diag}(f'(\mathbf{z}))$$

# 其它雅可比矩阵

- $\frac{\partial}{\partial \mathbf{x}} (\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{W}$
- $\frac{\partial}{\partial \mathbf{b}} (\mathbf{W}\mathbf{x} + \mathbf{b}) = \mathbf{I}$  (单位矩阵)
- $\frac{\partial}{\partial \mathbf{u}} (\mathbf{u}^T \mathbf{h}) = \mathbf{h}^T$

# 计算score的简单神经网络

- $s = \mathbf{u}^T \mathbf{h}$
- $\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$
- $\mathbf{x}$  (输入)



- 如何求  $\frac{\partial s}{\partial \mathbf{b}}$  ?

# 1. 把等式分解

- $s = \mathbf{u}^T \mathbf{h}$

- $\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b}) \rightarrow \mathbf{h} = f(\mathbf{z})$   
 $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$

- $\mathbf{x}$  (输入)

## 2. 应用链式规则

- $s = \mathbf{u}^T \mathbf{h}$

- $\mathbf{h} = f(\mathbf{z})$

- $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$

- $\mathbf{x}$  (输入)

- $$\begin{aligned} \frac{\partial s}{\partial \mathbf{b}} &= \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}} = \mathbf{u}^T \circ \text{diag}(f'(\mathbf{z})) \mathbf{I} \\ &= \mathbf{u}^T \circ \text{diag}(f'(\mathbf{z})) \end{aligned}$$



# 计算重用

- 如果我们想计算  $\frac{\partial s}{\partial \mathbf{W}}$
- 应用链式规则

$$\frac{\partial s}{\partial \mathbf{W}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$$

$$\frac{\partial s}{\partial \mathbf{b}} = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{b}}$$

- 令  $\delta = \frac{\partial s}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{z}}$ , 则  $\frac{\partial s}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$

# 对矩阵求导

- $\frac{\partial s}{\partial \mathbf{W}}$  的结果是什么形状?  $\mathbf{W} \in \mathbb{R}^{n \times m}$
- 1 output, nm inputs: 1 by nm Jacobian?
- 不方便进行参数更新  $\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$
- 解决办法: 导数的形状(shape)于参数的形状保持一致
- $\frac{\partial s}{\partial \mathbf{W}}$  的结果被表示成  $n \times m$  的矩阵 
$$\begin{bmatrix} \frac{\partial s}{\partial W_{11}} & \cdots & \frac{\partial s}{\partial W_{1m}} \\ \vdots & \ddots & \vdots \\ \frac{\partial s}{\partial W_{n1}} & \cdots & \frac{\partial s}{\partial W_{nm}} \end{bmatrix}$$

# 对矩阵求导

- 前面我们已经得到  $\frac{\partial s}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$ 
  - $\delta$  的计算前面已经完成
  - $\frac{\partial \mathbf{z}}{\partial \mathbf{W}}$  应该是  $\mathbf{x}$      $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$
- $\frac{\partial s}{\partial \mathbf{W}}$  的结果可以写成  $\delta^T \mathbf{x}^T$

$\delta$  is local error signal at  $\mathbf{z}$

$\mathbf{x}$  is local input signal

为什么转置?

$$\frac{\partial s}{\partial \mathbf{W}} = \boldsymbol{\delta}^T \mathbf{x}^T$$
$$[n \times m] \quad [n \times 1][1 \times m]$$


$$\frac{\partial s}{\partial \mathbf{W}} = \boldsymbol{\delta}^T \mathbf{x}^T = \begin{bmatrix} \delta_1 \\ \vdots \\ \delta_n \end{bmatrix} [x_1, \dots, x_m] = \begin{bmatrix} \delta_1 x_1 & \dots & \delta_1 x_m \\ \vdots & \ddots & \vdots \\ \delta_n x_1 & \dots & \delta_n x_m \end{bmatrix}$$

# 导数应该是什么形状？

- Disagreement between Jacobian form (which makes the chain rule easy) and the shape convention (which makes implementing SGD easy)
  - We expect answers to follow the **shape convention**
  - But Jacobian form is useful for computing the answers
- Two options:
  - 1. Use Jacobian form as much as possible, reshape to follow the convention at the end:
  - 2. Always follow the convention
    - Look at dimensions to figure out when to transpose and/or reorder terms.


# 面向反向传播的求导

- $\frac{\partial s}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \delta \frac{\partial}{\partial \mathbf{W}} \mathbf{W}\mathbf{x} + \mathbf{b}$
- 如何对 $\mathbf{W}$ 中的每一个权重 $W_{ij}$ 求导?
- $W_{ij}$  只用于计算  $z_i$ 
  - 例如 $W_{23}$  只用于计算  $z_2$
- $$\frac{\partial z_i}{\partial W_{ij}} = \frac{\partial}{\partial W_{ij}} \mathbf{W}_i \mathbf{x} + b_i$$
$$= \frac{\partial}{\partial W_{ij}} \sum_{k=1} W_{ik} x_k + b_i = x_j$$



# 面向反向传播的求导

- 所以得分 $s$ 对单个 $W_{ij}$ 求导的结果为

- $\frac{\partial s}{\partial W_{ij}} = \delta_i x_j$ 

Error signal from above

Local gradient signal

- 因此，对于整个矩阵 $\mathbf{W}$ 的求导结果可以写成

$$\frac{\partial s}{\partial \mathbf{W}} = \boldsymbol{\delta}^T \mathbf{x}^T$$
$$[n \times m] \quad [n \times 1][1 \times m]$$

# 求导过程中的小提示

- Carefully define your variables and keep track of their dimensionality!
- Chain rule! If  $y = f(u)$  and  $u = g(x)$ , i.e.,  $y = f(g(x))$ , then:
$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial u} \frac{\partial u}{\partial x}$$
- For the top softmax part of a model: First consider the derivative wrt  $f_c$  when  $c = y$  (the correct class), then consider derivative wrt  $f_c$  when  $c \neq y$  (all the incorrect classes)
- Work out element-wise partial derivatives if you're getting confused by matrix calculus!
- Use Shape Convention. Note: The error message  $\delta$  that arrives at a hidden layer has the same dimensionality as that hidden layer



# 主要内容

- 线性分类方法与神经网络
- 神经元与神经网络的表示
- 神经网络在NLP中的应用示例
- 神经网络中导数的计算
- 计算图与反向传播

# 计算图与反向传播 (Computation Graphs and Backpropagation)

- 我们可以用一个图表示前面提出的神经网络


- Source nodes: inputs
- Interior nodes: operations

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (输入)}$$



# 计算图与反向传播 (Computation Graphs and Backpropagation)

- 我们可以用一个图表示前面提出的神经网络

- Source nodes: inputs
- Interior nodes: operations
- Edges pass along result of the operation


$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (输入)}$$

“Forward Propagation”



# 反向传播 (Backpropagation)

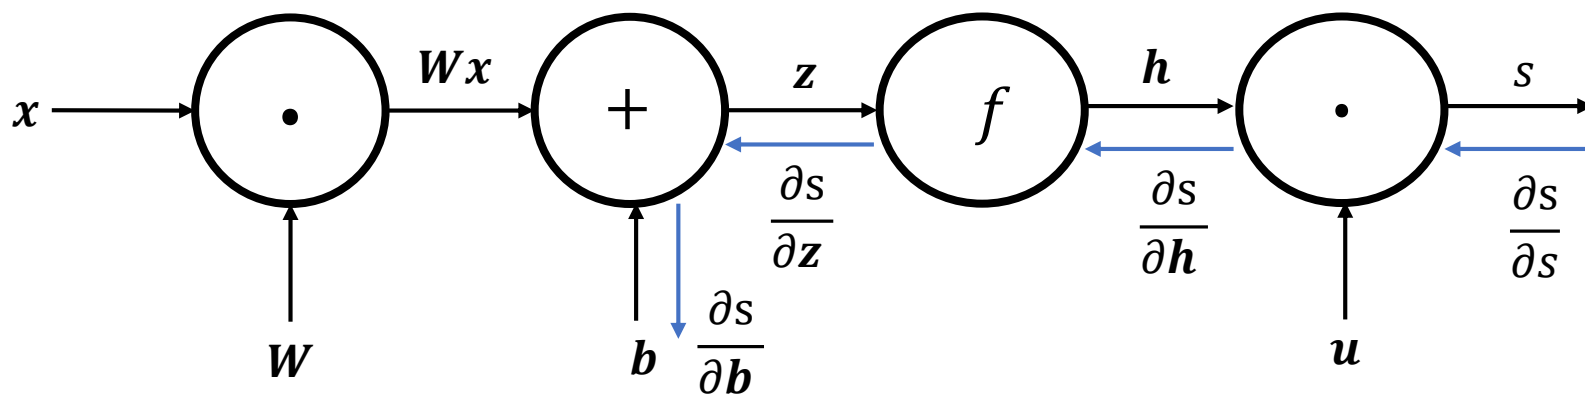
- *Go backwards along edges*
  - Pass along **gradients**

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$


$\mathbf{x}$  (输入)



# 反向传播中的单个节点

- Node receives an “upstream gradient”
- Goal is to pass on the correct “downstream gradient”


$$h = f(z)$$



# 反向传播中的单个节点

- Each node has a **local gradient**
  - The gradient of it's output with respect to it's input

$$h = f(z)$$




# 反向传播中的单个节点

- Each node has a **local gradient**

- The gradient of it's output with respect to it's input

$$h = f(z)$$


- [downstream gradient] = [upstream gradient] x [local gradient]



# 反向传播中的单个节点

- 具有多个输入的节点该如何处理？

$$z = Wx$$






# 反向传播中的单个节点

- 具有多个输入的节点该如何处理？

$$z = Wx$$



Downstream  
gradient

Local  
gradient

Upstream  
gradient

# 一个简单示例

- 前向传播步骤


$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$



# 一个简单示例

- 前向传播步骤


$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$



# 一个简单示例

- 前向传播步骤


$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 1, y = 2, z = 0$$



# 一个简单示例

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- 前向传播步骤

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



# 一个简单示例

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- 前向传播步骤

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



upstream \* local = downstream

# 一个简单示例

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- 前向传播步骤

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



upstream \* local = downstream

# 一个简单示例

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- 前向传播步骤

$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$

$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$



upstream \* local = downstream



# 一个简单示例

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- 前向传播步骤

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$

Local gradients

$$\frac{\partial a}{\partial x} = 1 \quad \frac{\partial a}{\partial y} = 1$$


$$\frac{\partial b}{\partial y} = \mathbf{1}(y > z) = 1 \quad \frac{\partial b}{\partial z} = \mathbf{1}(z > y) = 0$$

$$\frac{\partial f}{\partial a} = b = 2 \quad \frac{\partial f}{\partial b} = a = 3$$

$$\frac{\partial f}{\partial x} = 2$$


$$\frac{\partial f}{\partial y} = 3 + 2 = 5$$

$$\frac{\partial f}{\partial z} = 0$$



upstream \* local = downstream

# Gradients sum at outward branches



$$a = x + y$$

$$b = \max(y, z)$$


$$f = ab$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial a} \frac{\partial a}{\partial y} + \frac{\partial f}{\partial b} \frac{\partial b}{\partial y}$$

# 对于节点的直观理解

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$


- + “distributes” the upstream gradient



# 对于节点的直观理解

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$


- + “distributes” the upstream gradient
- max “routes” the upstream gradient



# 对于节点的直观理解

$$f(x, y, z) = (x + y) \max(y, z)$$
$$x = 1, y = 2, z = 0$$

- + “distributes” the upstream gradient
- max “routes” the upstream gradient
- \* “switches” the upstream gradient



# 反向传播算法的效率

- Incorrect way of doing backprop

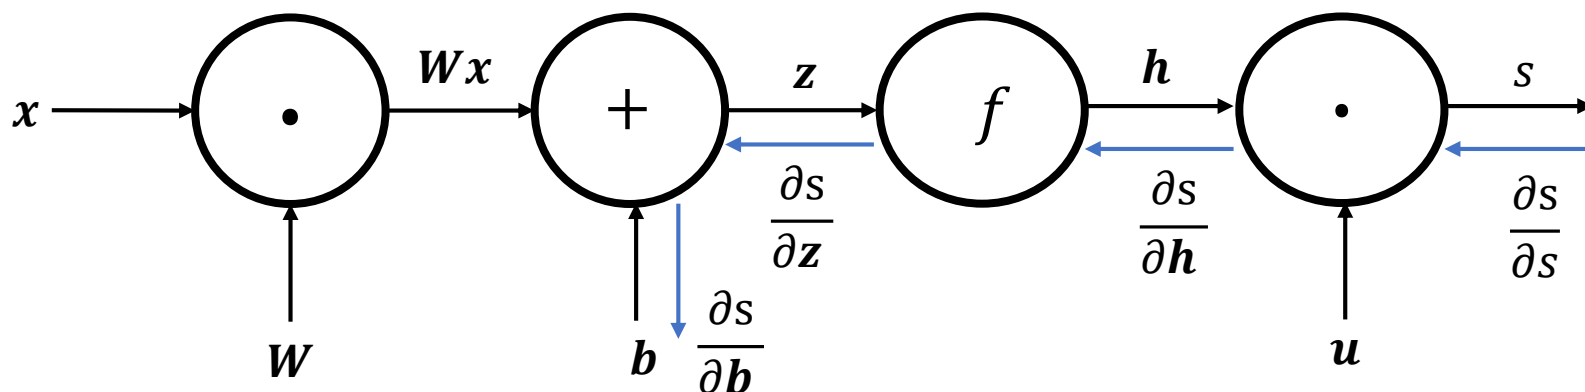
- 先计算  $\frac{\partial s}{\partial b}$

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{x} \text{ (输入)}$$



# 反向传播算法的效率

- Incorrect way of doing backprop


- 先计算  $\frac{\partial s}{\partial b}$
- 再独立计算  $\frac{\partial s}{\partial W}$
- *Duplicated computation!*

$$s = \mathbf{u}^T \mathbf{h}$$

$$\mathbf{h} = f(\mathbf{z})$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$


$$\mathbf{x} \text{ (输入)}$$



# 反向传播算法的效率


- Correct way
  - *Compute all the gradients at once*
  - *Analogous to using  $\delta$  when we computed gradients by hand*

$$\begin{aligned}s &= \mathbf{u}^T \mathbf{h} \\ \mathbf{h} &= f(\mathbf{z}) \\ \mathbf{z} &= \mathbf{W}\mathbf{x} + \mathbf{b} \\ \mathbf{x} &(\text{输入})\end{aligned}$$





# Back-Prop in General Computation Graph



1. Fprop: visit nodes in topological sort order
  - Compute value of node given predecessors
2. Bprop:
  - initialize output gradient = 1
  - visit nodes in reverse order:



- Compute gradient wrt each node using gradient wrt successors
- $\{y_1, y_2, \dots, y_n\} = \text{successors of } x$

$$\frac{\partial z}{\partial x} = \sum_{i=1}^n \frac{\partial z}{\partial y_i} \frac{\partial y_i}{\partial x}$$

Done correctly, big  $O()$  complexity of fprop and bprop is **the same**

In general our nets have regular layer-structure and so we can use matrices and Jacobians...

# Automatic Differentiation




- The gradient computation can be automatically inferred from the symbolic expression of the fprop
- Each node type needs to know how to compute its output and how to compute the gradient wrt its inputs given the gradient wrt its output
- Modern DL frameworks (Tensorflow, PyTorch, etc.) do backpropagation for you but mainly leave layer/node writer to hand-calculate the local derivative

# Backprop Implementations

```
class ComputationalGraph(object):  
    #...  
    def forward(inputs):  
        # 1. [pass inputs to input gates...]  
        # 2. forward the computational graph:  
        for gate in self.graph.nodes_topologically_sorted():  
            gate.forward()  
        return loss # the final gate in the graph outputs the loss  
    def backward():  
        for gate in reversed(self.graph.nodes_topologically_sorted()):  
            gate.backward() # little piece of backprop (chain rule applied)  
        return inputs_gradients
```

# Implementation: forward/backward API




(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        return z  
    def backward(dz):  
        # dx = ... #todo  
        # dy = ... #todo  
        return [dx, dy]
```

$$\frac{\partial L}{\partial z}$$

$$\frac{\partial L}{\partial x}$$

# Implementation: forward/backward API



(x,y,z are scalars)

```
class MultiplyGate(object):  
    def forward(x,y):  
        z = x*y  
        self.x = x # must keep these around!  
        self.y = y  
        return z  
    def backward(dz):  
        dx = self.y * dz # [dz/dx * dL/dz]  
        dy = self.x * dz # [dz/dy * dL/dz]  
        return [dx, dy]
```

# 总结

- 神经网络是一种非线性分类方法
- 神经网络参数的优化可以采用随机梯度下降 (SGD) 的方法
- SGD中的参数更新需要用到反向传播方法 (Backpropagation)
- 反向传播: recursively apply the chain rule along computation graph
  - $[\text{downstream gradient}] = [\text{upstream gradient}] \times [\text{local gradient}]$
  - Forward pass: compute results of operations and save intermediate values
  - Backward pass: apply chain rule to compute gradients